

# Tanawin-st123975-vgg16

October 20, 2023

```
[ ]: !nvidia-smi
```

Thu Oct 19 21:01:56 2023

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		525.105.17		Driver Version: 525.105.17			CUDA Version: 12.0		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan Temp Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M.	
								MIG M.	
+=====+=====+=====+=====+=====+=====+=====+=====+=====+=====+									
0 Tesla T4		Off		00000000:00:04.0 Off				0	
N/A 57C P0		29W / 70W		3263MiB / 15360MiB		0%		Default	
								N/A	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU GI CI		PID		Type		Process name		GPU Memory	
ID ID								Usage	
+=====+=====+=====+=====+=====+=====+=====+=====+=====+=====+									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

```
[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import seaborn as sns
import numpy as np
from torch.utils.data import random_split
device=torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[ ]: transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
```

```

])

# transform_train = transforms.Compose([
#     transforms.RandomResizedCrop(224),
#     transforms.RandomHorizontalFlip(),
#     transforms.ToTensor(),
#     transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
# ↪225])
# ])

# Normalization for testing, no data augmentation
transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

# transform_train = transforms.Compose([
#     transforms.RandomResizedCrop(32, scale=(0.8, 1.0)),
#     transforms.RandomHorizontalFlip(),
#     transforms.ColorJitter(brightness=0.4, contrast=0.4, saturation=0.4,
# ↪hue=0.1),
#     transforms.RandomRotation(20),
#     transforms.ToTensor(),
#     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
# ])

# transform_test = transforms.Compose([
#     transforms.ToTensor(),
#     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
# ])

batch_size = 8

# Load CIFAR-10 dataset
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True,
                                         ↪transform=transform_train)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                         shuffle=True, num_workers=32)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=batch_size,
                                         shuffle=False, num_workers=32)

```

```
classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:560:
UserWarning: This DataLoader will create 32 worker processes in total. Our
suggested max number of worker in current system is 8, which is smaller than
what this DataLoader is going to create. Please be aware that excessive worker
creation might get DataLoader running slow or even freeze, lower the worker
number to avoid potential slowness/freeze if necessary.
```

```
warnings.warn(_create_warning_msg(
```

Files already downloaded and verified


```
[ ]: import matplotlib.pyplot as plt
import numpy as np

# functions to show an image

def imshow(img):
    img = img / 2 + 0.5     # unnormalize
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()

# get some random training images
dataiter = iter(trainloader)
images, labels = next(dataiter)

# show images
imshow(torchvision.utils.make_grid(images))
# print labels
print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))
```



```
truck dog    truck bird   cat    truck truck horse
```

```
[ ]: # import torch
# import torch.nn as nn
# import torch.nn.functional as F

# class VGG16_NET(nn.Module):
#     def __init__(self):
#         super(VGG16_NET, self).__init__()
#         self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
# ↪padding=1)
#         self.bn1 = nn.BatchNorm2d(64)
#         self.conv2 = nn.Conv2d(in_channels=64, out_channels=64,
# ↪kernel_size=3, padding=1)
#         self.bn2 = nn.BatchNorm2d(64)

#         self.conv3 = nn.Conv2d(in_channels=64, out_channels=128,
# ↪kernel_size=3, padding=1)
#         self.bn3 = nn.BatchNorm2d(128)
#         self.conv4 = nn.Conv2d(in_channels=128, out_channels=128,
# ↪kernel_size=3, padding=1)
#         self.bn4 = nn.BatchNorm2d(128)

#         self.conv5 = nn.Conv2d(in_channels=128, out_channels=256,
# ↪kernel_size=3, padding=1)
#         self.bn5 = nn.BatchNorm2d(256)
#         self.conv6 = nn.Conv2d(in_channels=256, out_channels=256,
# ↪kernel_size=3, padding=1)
#         self.bn6 = nn.BatchNorm2d(256)
#         self.conv7 = nn.Conv2d(in_channels=256, out_channels=256,
# ↪kernel_size=3, padding=1)
#         self.bn7 = nn.BatchNorm2d(256)

#         self.conv8 = nn.Conv2d(in_channels=256, out_channels=512,
# ↪kernel_size=3, padding=1)
#         self.bn8 = nn.BatchNorm2d(512)
#         self.conv9 = nn.Conv2d(in_channels=512, out_channels=512,
# ↪kernel_size=3, padding=1)
#         self.bn9 = nn.BatchNorm2d(512)
#         self.conv10 = nn.Conv2d(in_channels=512, out_channels=512,
# ↪kernel_size=3, padding=1)
#         self.bn10 = nn.BatchNorm2d(512)

#         self.conv11 = nn.Conv2d(in_channels=512, out_channels=512,
# ↪kernel_size=3, padding=1)
#         self.bn11 = nn.BatchNorm2d(512)
#         self.conv12 = nn.Conv2d(in_channels=512, out_channels=512,
# ↪kernel_size=3, padding=1)
```

```

#         self.bn12 = nn.BatchNorm2d(512)
#         self.conv13 = nn.Conv2d(in_channels=512, out_channels=512,
↪kernel_size=3, padding=1)
#         self.bn13 = nn.BatchNorm2d(512)

#         self.maxpool = nn.MaxPool2d(kernel_size=2, stride=2)

#         self.fc14 = nn.Linear(512, 4096)
#         self.fc15 = nn.Linear(4096, 4096)
#         self.fc16 = nn.Linear(4096, 10)

#     def forward(self, x):
#         x = F.relu(self.bn1(self.conv1(x)))
#         x = F.relu(self.bn2(self.conv2(x)))
#         x = self.maxpool(x)
#         x = F.relu(self.bn3(self.conv3(x)))
#         x = F.relu(self.bn4(self.conv4(x)))
#         x = self.maxpool(x)
#         x = F.relu(self.bn5(self.conv5(x)))
#         x = F.relu(self.bn6(self.conv6(x)))
#         x = F.relu(self.bn7(self.conv7(x)))
#         x = self.maxpool(x)
#         x = F.relu(self.bn8(self.conv8(x)))
#         x = F.relu(self.bn9(self.conv9(x)))
#         x = F.relu(self.bn10(self.conv10(x)))
#         x = self.maxpool(x)
#         x = F.relu(self.bn11(self.conv11(x)))
#         x = F.relu(self.bn12(self.conv12(x)))
#         x = F.relu(self.bn13(self.conv13(x)))
#         x = self.maxpool(x)

#         x = x.view(x.size(0), -1)
#         x = F.relu(self.fc14(x))
#         x = F.dropout(x, 0.5)
#         x = F.relu(self.fc15(x))
#         x = F.dropout(x, 0.5)
#         x = self.fc16(x)
#         return x

# # Create an instance of the combined model
# device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
# model = VGG16_NET()
# model.to(device)

```

```

[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F

```

```

class VGG16_NET(nn.Module):
    def __init__(self):
        super(VGG16_NET, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=3, out_channels=64, kernel_size=3,
        ↪padding=1)
        self.bn1 = nn.BatchNorm2d(64)
        self.conv2 = nn.Conv2d(in_channels=64, out_channels=64, kernel_size=3,
        ↪padding=1)
        self.bn2 = nn.BatchNorm2d(64)

        self.conv3 = nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3,
        ↪padding=1)
        self.bn3 = nn.BatchNorm2d(128)
        self.conv4 = nn.Conv2d(in_channels=128, out_channels=128,
        ↪kernel_size=3, padding=1)
        self.bn4 = nn.BatchNorm2d(128)

        self.conv5 = nn.Conv2d(in_channels=128, out_channels=256,
        ↪kernel_size=3, padding=1)
        self.bn5 = nn.BatchNorm2d(256)
        self.conv6 = nn.Conv2d(in_channels=256, out_channels=256,
        ↪kernel_size=3, padding=1)
        self.bn6 = nn.BatchNorm2d(256)
        self.conv7 = nn.Conv2d(in_channels=256, out_channels=256,
        ↪kernel_size=3, padding=1)
        self.bn7 = nn.BatchNorm2d(256)

        self.conv8 = nn.Conv2d(in_channels=256, out_channels=512,
        ↪kernel_size=3, padding=1)
        self.bn8 = nn.BatchNorm2d(512)
        self.conv9 = nn.Conv2d(in_channels=512, out_channels=512,
        ↪kernel_size=3, padding=1)
        self.bn9 = nn.BatchNorm2d(512)
        self.conv10 = nn.Conv2d(in_channels=512, out_channels=512,
        ↪kernel_size=3, padding=1)
        self.bn10 = nn.BatchNorm2d(512)

        self.conv11 = nn.Conv2d(in_channels=512, out_channels=512,
        ↪kernel_size=3, padding=1)
        self.bn11 = nn.BatchNorm2d(512)
        self.conv12 = nn.Conv2d(in_channels=512, out_channels=512,
        ↪kernel_size=3, padding=1)
        self.bn12 = nn.BatchNorm2d(512)
        self.conv13 = nn.Conv2d(in_channels=512, out_channels=512,
        ↪kernel_size=3, padding=1)

```

```

self.bn13 = nn.BatchNorm2d(512)

self.maxpool = nn.MaxPool2d(kernel_size=2, stride=2)

self.fc14 = nn.Linear(512, 4096)
self.dropout1 = nn.Dropout(0.5)
self.fc15 = nn.Linear(4096, 4096)
self.dropout2 = nn.Dropout(0.5)
self.fc16 = nn.Linear(4096, 10)

def forward(self, x):
    x = F.relu(self.bn1(self.conv1(x)))
    x = F.relu(self.bn2(self.conv2(x)))
    x = self.maxpool(x)
    x = F.relu(self.bn3(self.conv3(x)))
    x = F.relu(self.bn4(self.conv4(x)))
    x = self.maxpool(x)
    x = F.relu(self.bn5(self.conv5(x)))
    x = F.relu(self.bn6(self.conv6(x)))
    x = F.relu(self.bn7(self.conv7(x)))
    x = self.maxpool(x)
    x = F.relu(self.bn8(self.conv8(x)))
    x = F.relu(self.bn9(self.conv9(x)))
    x = F.relu(self.bn10(self.conv10(x)))
    x = self.maxpool(x)
    x = F.relu(self.bn11(self.conv11(x)))
    x = F.relu(self.bn12(self.conv12(x)))
    x = F.relu(self.bn13(self.conv13(x)))
    x = self.maxpool(x)

    x = x.view(x.size(0), -1)
    x = F.relu(self.fc14(x))
    x = self.dropout1(x)
    x = F.relu(self.fc15(x))
    x = self.dropout2(x)
    x = self.fc16(x)
    return x

# Create an instance of the combined model
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = VGG16_NET()
model.to(device)

```

```

[ ]: VGG16_NET(
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv4): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn4): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv5): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn5): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv6): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv7): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn7): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv8): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn8): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv9): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv10): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn10): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv11): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn11): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv12): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn12): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (conv13): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (bn13): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (maxpool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (fc14): Linear(in_features=512, out_features=4096, bias=True)
    (dropout1): Dropout(p=0.5, inplace=False)
    (fc15): Linear(in_features=4096, out_features=4096, bias=True)
    (dropout2): Dropout(p=0.5, inplace=False)
    (fc16): Linear(in_features=4096, out_features=10, bias=True)
)

```

```
[ ]: # model._dropout = nn.Dropout(p=0.5)
```



```
[ ]: import torch.optim as optim

criterion = nn.CrossEntropyLoss()

# optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9,
# ↪weight_decay=0.0001)
# optimizer = optim.RMSprop(model.parameters(), lr=0.001, alpha=0.9,
# ↪weight_decay=0.0001)
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=0.0001)
```

```
[ ]: from torch.optim.lr_scheduler import ExponentialLR, MultiStepLR

scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.1)

scheduler1 = ExponentialLR(optimizer, gamma=0.9)
scheduler2 = MultiStepLR(optimizer, milestones=[30,80], gamma=0.1)
```

```
[ ]: model_name = "vgg16_scratch"
# Training loop
num_epochs = 50
train_loss_list = []
valid_loss_list = []
train_accuracy_list = []
valid_accuracy_list = []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in trainloader:
        images = images.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
```

```

train_loss = running_loss / len(trainloader)
train_accuracy = correct / total

# Validation loop
model.eval()
running_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for images, labels in testloader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

valid_loss = running_loss / len(testloader)
valid_accuracy = correct / total

# Store loss and accuracy values
train_loss_list.append(train_loss)
valid_loss_list.append(valid_loss)
train_accuracy_list.append(train_accuracy)
valid_accuracy_list.append(valid_accuracy)

# Print the training/validation statistics
print(f"Epoch: {epoch+1}/{num_epochs} | "
      f"Train Loss: {train_loss:.4f} | Train Acc: {train_accuracy:.4f} | "
      f"Valid Loss: {valid_loss:.4f} | Valid Acc: {valid_accuracy:.4f}")

# Update the learning rate
# scheduler.step()
scheduler1.step()
scheduler2.step()

# Save the trained model
torch.save(model.state_dict(), model_name+".pth")

```

```

Epoch: 1/50 | Train Loss: 2.1251 | Train Acc: 0.1621 | Valid Loss: 1.9477 |
Valid Acc: 0.1917
Epoch: 2/50 | Train Loss: 1.9899 | Train Acc: 0.1898 | Valid Loss: 1.8724 |

```

Valid Acc: 0.2290  
Epoch: 3/50 | Train Loss: 1.9107 | Train Acc: 0.2286 | Valid Loss: 1.7506 |  
Valid Acc: 0.2818  
Epoch: 4/50 | Train Loss: 1.7977 | Train Acc: 0.2776 | Valid Loss: 1.5297 |  
Valid Acc: 0.3790  
Epoch: 5/50 | Train Loss: 1.5679 | Train Acc: 0.3777 | Valid Loss: 1.3371 |  
Valid Acc: 0.4682  
Epoch: 6/50 | Train Loss: 1.3651 | Train Acc: 0.4771 | Valid Loss: 1.1292 |  
Valid Acc: 0.5640  
Epoch: 7/50 | Train Loss: 1.1854 | Train Acc: 0.5696 | Valid Loss: 1.0025 |  
Valid Acc: 0.6371  
Epoch: 8/50 | Train Loss: 1.0637 | Train Acc: 0.6171 | Valid Loss: 0.8898 |  
Valid Acc: 0.6758  
Epoch: 9/50 | Train Loss: 0.9780 | Train Acc: 0.6542 | Valid Loss: 0.8506 |  
Valid Acc: 0.6928  
Epoch: 10/50 | Train Loss: 0.9052 | Train Acc: 0.6816 | Valid Loss: 0.8321 |  
Valid Acc: 0.7051  
Epoch: 11/50 | Train Loss: 0.8479 | Train Acc: 0.7052 | Valid Loss: 0.7546 |  
Valid Acc: 0.7368  
Epoch: 12/50 | Train Loss: 0.7956 | Train Acc: 0.7281 | Valid Loss: 0.6937 |  
Valid Acc: 0.7607  
Epoch: 13/50 | Train Loss: 0.7444 | Train Acc: 0.7454 | Valid Loss: 0.6799 |  
Valid Acc: 0.7693  
Epoch: 14/50 | Train Loss: 0.7087 | Train Acc: 0.7608 | Valid Loss: 0.6796 |  
Valid Acc: 0.7733  
Epoch: 15/50 | Train Loss: 0.6712 | Train Acc: 0.7746 | Valid Loss: 0.6472 |  
Valid Acc: 0.7776  
Epoch: 16/50 | Train Loss: 0.6456 | Train Acc: 0.7834 | Valid Loss: 0.6030 |  
Valid Acc: 0.7989  
Epoch: 17/50 | Train Loss: 0.6143 | Train Acc: 0.7931 | Valid Loss: 0.5960 |  
Valid Acc: 0.7996  
Epoch: 18/50 | Train Loss: 0.5920 | Train Acc: 0.8027 | Valid Loss: 0.5684 |  
Valid Acc: 0.8123  
Epoch: 19/50 | Train Loss: 0.5678 | Train Acc: 0.8115 | Valid Loss: 0.5836 |  
Valid Acc: 0.8068  
Epoch: 20/50 | Train Loss: 0.5457 | Train Acc: 0.8181 | Valid Loss: 0.5346 |  
Valid Acc: 0.8215  
Epoch: 21/50 | Train Loss: 0.5244 | Train Acc: 0.8258 | Valid Loss: 0.5224 |  
Valid Acc: 0.8253  
Epoch: 22/50 | Train Loss: 0.5073 | Train Acc: 0.8314 | Valid Loss: 0.5436 |  
Valid Acc: 0.8190  
Epoch: 23/50 | Train Loss: 0.4969 | Train Acc: 0.8357 | Valid Loss: 0.5074 |  
Valid Acc: 0.8301  
Epoch: 24/50 | Train Loss: 0.4781 | Train Acc: 0.8381 | Valid Loss: 0.5102 |  
Valid Acc: 0.8284  
Epoch: 25/50 | Train Loss: 0.4696 | Train Acc: 0.8419 | Valid Loss: 0.4934 |  
Valid Acc: 0.8406  
Epoch: 26/50 | Train Loss: 0.4563 | Train Acc: 0.8472 | Valid Loss: 0.4893 |


Valid Acc: 0.8377  
Epoch: 27/50 | Train Loss: 0.4483 | Train Acc: 0.8505 | Valid Loss: 0.4887 |  
Valid Acc: 0.8386  
Epoch: 28/50 | Train Loss: 0.4380 | Train Acc: 0.8547 | Valid Loss: 0.4899 |  
Valid Acc: 0.8398  
Epoch: 29/50 | Train Loss: 0.4296 | Train Acc: 0.8570 | Valid Loss: 0.4844 |  
Valid Acc: 0.8425  
Epoch: 30/50 | Train Loss: 0.4222 | Train Acc: 0.8576 | Valid Loss: 0.4867 |  
Valid Acc: 0.8399  
Epoch: 31/50 | Train Loss: 0.4013 | Train Acc: 0.8667 | Valid Loss: 0.4783 |  
Valid Acc: 0.8434  
Epoch: 32/50 | Train Loss: 0.4000 | Train Acc: 0.8671 | Valid Loss: 0.4704 |  
Valid Acc: 0.8476  
Epoch: 33/50 | Train Loss: 0.3948 | Train Acc: 0.8687 | Valid Loss: 0.4695 |  
Valid Acc: 0.8468  
Epoch: 34/50 | Train Loss: 0.3943 | Train Acc: 0.8683 | Valid Loss: 0.4681 |  
Valid Acc: 0.8470  
Epoch: 35/50 | Train Loss: 0.3935 | Train Acc: 0.8679 | Valid Loss: 0.4691 |  
Valid Acc: 0.8472  
Epoch: 36/50 | Train Loss: 0.3917 | Train Acc: 0.8687 | Valid Loss: 0.4685 |  
Valid Acc: 0.8464  
Epoch: 37/50 | Train Loss: 0.3945 | Train Acc: 0.8701 | Valid Loss: 0.4665 |  
Valid Acc: 0.8466  
Epoch: 38/50 | Train Loss: 0.3905 | Train Acc: 0.8706 | Valid Loss: 0.4703 |  
Valid Acc: 0.8453  
Epoch: 39/50 | Train Loss: 0.3931 | Train Acc: 0.8691 | Valid Loss: 0.4689 |  
Valid Acc: 0.8483  
Epoch: 40/50 | Train Loss: 0.3934 | Train Acc: 0.8690 | Valid Loss: 0.4646 |  
Valid Acc: 0.8476  
Epoch: 41/50 | Train Loss: 0.3911 | Train Acc: 0.8697 | Valid Loss: 0.4662 |  
Valid Acc: 0.8463  
Epoch: 42/50 | Train Loss: 0.3877 | Train Acc: 0.8702 | Valid Loss: 0.4686 |  
Valid Acc: 0.8465  
Epoch: 43/50 | Train Loss: 0.3869 | Train Acc: 0.8704 | Valid Loss: 0.4694 |  
Valid Acc: 0.8468  
Epoch: 44/50 | Train Loss: 0.3886 | Train Acc: 0.8702 | Valid Loss: 0.4647 |  
Valid Acc: 0.8473  
Epoch: 45/50 | Train Loss: 0.3842 | Train Acc: 0.8716 | Valid Loss: 0.4653 |  
Valid Acc: 0.8474  
Epoch: 46/50 | Train Loss: 0.3899 | Train Acc: 0.8687 | Valid Loss: 0.4638 |  
Valid Acc: 0.8479  
Epoch: 47/50 | Train Loss: 0.3882 | Train Acc: 0.8702 | Valid Loss: 0.4713 |  
Valid Acc: 0.8463  
Epoch: 48/50 | Train Loss: 0.3877 | Train Acc: 0.8714 | Valid Loss: 0.4700 |  
Valid Acc: 0.8467  
Epoch: 49/50 | Train Loss: 0.3836 | Train Acc: 0.8722 | Valid Loss: 0.4670 |  
Valid Acc: 0.8464  
Epoch: 50/50 | Train Loss: 0.3875 | Train Acc: 0.8711 | Valid Loss: 0.4675 |

Valid Acc: 0.8482

```
[ ]: import matplotlib.pyplot as plt
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
# Plot training and validation loss
plt.plot(range(1, num_epochs+1), train_loss_list, label='Train Loss')
plt.plot(range(1, num_epochs+1), valid_loss_list, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

plt.subplot(1, 2, 2)
# Plot training and validation accuracy
plt.plot(range(1, num_epochs+1), train_accuracy_list, label='Train Accuracy')
plt.plot(range(1, num_epochs+1), valid_accuracy_list, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.suptitle(model_name)
plt.show()
```



```
[ ]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import numpy as np
```

```

loaded_model = VGG16_NET()
loaded_model.load_state_dict(torch.load('/content/vgg16_scratch.pth'))
loaded_model.to(device) # Move the loaded model to the same device as the data

# Test the loaded model
def test_loaded_model(model, test_dataloader, device):
    model.eval() # Set the model to evaluation mode
    all_preds = []
    all_labels = []
    with torch.no_grad():
        for inputs, labels in test_dataloader:
            inputs = inputs.to(device)
            labels = labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)
            all_preds.extend(preds.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
    return all_preds, all_labels

# Test the loaded model
all_preds, all_labels = test_loaded_model(loaded_model, testloader, device)

# Generate confusion matrix
cm = confusion_matrix(all_labels, all_preds)

# Display the confusion matrix using ConfusionMatrixDisplay
cmd = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.
    ↪unique(all_labels))
cmd.plot(cmap='Blues', xticks_rotation='vertical')

# Show the plot
plt.show()

```

