

Tanawin-st123975-ViT-v100

November 17, 2023

```
[1]: !nvidia-smi
```

Fri Nov 17 07:29:33 2023

```
+-----+
| NVIDIA-SMI 525.105.17      Driver Version: 525.105.17      CUDA Version: 12.0      |
+-----+-----+-----+-----+-----+
| GPU  Name                Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|                                       |                    |    MIG M.     |
+-----+-----+-----+-----+-----+
|   0   Tesla V100-SXM2...    Off      | 00000000:00:04:0  Off  |           0         |
| N/A   34C    P0      24W / 300W |      0MiB / 16384MiB |           0%      Default |
|                                       |                    |           N/A     |
+-----+-----+-----+-----+-----+

+-----+
| Processes:                                     |
|  GPU   GI    CI          PID    Type    Process name                        GPU Memory |
|          ID    ID                                   Usage                        |
+-----+-----+-----+-----+-----+
| No running processes found                                     |
+-----+
```

```
[2]: # from google.colab import files
      # files.upload()
```

```
[3]: # import os

      # # Set Kaggle API credentials as environment variables
      # os.environ['KAGGLE_USERNAME'] = 'kmutnb'
      # os.environ['KAGGLE_KEY'] = '2e8f3eec48ac90959791a330f9109431'

      # import kaggle

      # # Authenticate with Kaggle API
      # kaggle.api.authenticate()
```

```
[4]: # !kaggle datasets download -d gpiosenka/sports-classification
```

```
[5]: # !unzip /content/sports-classification.zip -d data
```

```
[6]: !pip install vit-pytorch linformer
```

Collecting vit-pytorch

Downloading vit_pytorch-1.6.4-py3-none-any.whl (100 kB)

100.3/100.3

kB 2.5 MB/s eta 0:00:00

Collecting linformer

Downloading linformer-0.2.1-py3-none-any.whl (6.1 kB)

Collecting einops>=0.7.0 (from vit-pytorch)

Downloading einops-0.7.0-py3-none-any.whl (44 kB)

44.6/44.6 kB

6.5 MB/s eta 0:00:00

Requirement already satisfied: torch>=1.10 in

/usr/local/lib/python3.10/dist-packages (from vit-pytorch) (2.1.0+cu118)

Requirement already satisfied: torchvision in /usr/local/lib/python3.10/dist-packages (from vit-pytorch) (0.16.0+cu118)

Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (3.13.1)

Requirement already satisfied: typing-extensions in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (4.5.0)

Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (1.12)

Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (3.2.1)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (3.1.2)

Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (2023.6.0)

Requirement already satisfied: triton==2.1.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.10->vit-pytorch) (2.1.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from torchvision->vit-pytorch) (1.23.5)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from torchvision->vit-pytorch) (2.31.0)

Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.10/dist-packages (from torchvision->vit-pytorch) (9.4.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.10->vit-pytorch) (2.1.3)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->torchvision->vit-pytorch) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->torchvision->vit-pytorch) (3.4)

Requirement already satisfied: urllib3<3,>=1.21.1 in

```
/usr/local/lib/python3.10/dist-packages (from requests->torchvision->vit-  
pytorch) (2.0.7)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests->torchvision->vit-  
pytorch) (2023.7.22)  
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-  
packages (from sympy->torch>=1.10->vit-pytorch) (1.3.0)  
Installing collected packages: einops, linformer, vit-pytorch  
Successfully installed einops-0.7.0 linformer-0.2.1 vit-pytorch-1.6.4
```

```
[7]: import glob  
from itertools import chain  
import os  
import random  
import zipfile  
  
import matplotlib.pyplot as plt  
import numpy as np  
import pandas as pd  
import torch  
import torch.nn as nn  
import torch.nn.functional as F  
import torch.optim as optim  
from linformer import Linformer  
from PIL import Image  
from sklearn.model_selection import train_test_split  
from torch.optim.lr_scheduler import StepLR  
from torch.utils.data import DataLoader, Dataset  
from torchvision import datasets, transforms  
from tqdm.notebook import tqdm  
  
from vit_pytorch.efficient import ViT  
  
import torch  
import torch.nn as nn  
import torch.optim as optim  
from torch.utils.data import DataLoader  
from torchvision.datasets import ImageFolder  
from torchvision.transforms import transforms
```

```
[8]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")  
device
```

```
[8]: device(type='cuda')
```

```
[9]: transform = transforms.Compose([  
    transforms.RandomResizedCrop(224),
```

```

        transforms.RandomHorizontalFlip(),
        transforms.RandomVerticalFlip(),
        transforms.ColorJitter(brightness=0.4, contrast=0.4, saturation=0.4, hue=0.
↪1),
        transforms.RandomRotation(30),
        transforms.RandomAffine(degrees=0, translate=(0.1, 0.1), scale=(0.9, 1.1),
↪shear=10),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

```

```

[10]: train_data = "/content/data/train"
      val_data = "/content/data/valid"

```

```

[11]: train_dataset = ImageFolder(train_data, transform=transform)
      valid_dataset = ImageFolder(val_data, transform=transform)

```

```

[12]: batch_size = 8
      # Define the data loaders
      train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True,
↪num_workers=4)
      valid_loader = DataLoader(valid_dataset, batch_size=batch_size, shuffle=False,
↪num_workers=4)

```

```

[13]: from vit_pytorch import ViT

      v = ViT(
          image_size = 256,
          patch_size = 32,
          num_classes = 100,
          dim = 1024,
          depth = 6,
          heads = 16,
          mlp_dim = 2048,
          dropout = 0.1,
          emb_dropout = 0.1
      )

```

```

[14]: from torchvision.models import vit_b_16 as ViT, ViT_B_16_Weights
      model = ViT(weights=ViT_B_16_Weights.DEFAULT)

      # from torchvision.models import vit_l_32 as ViT, ViT_L_32_Weights
      # model = ViT(weights=ViT_L_32_Weights.DEFAULT)

      model = model.to(device)

```

Downloading: "https://download.pytorch.org/models/vit_b_16-c867db91.pth" to

```
/root/.cache/torch/hub/checkpoints/vit_b_16-c867db91.pth
100%|      | 330M/330M [00:05<00:00, 67.6MB/s]
```

```
[15]: model._dropout = nn.Dropout(p=0.5)
```

```
[16]: # Define the loss function and optimizer
criterion = nn.CrossEntropyLoss()
# optimizer = optim.Adam(model.parameters(), lr=0.001)
# optimizer = optim.Adam(model.parameters(), lr=0.000001)
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
```

```
[17]: # Define the learning rate scheduler
scheduler = optim.lr_scheduler.StepLR(optimizer, step_size=5, gamma=0.1)
```

```
[18]: model_name = "ViT-"+str(batch_size)+"-15-nov-23"
# Training loop
num_epochs = 10
train_loss_list = []
valid_loss_list = []
train_accuracy_list = []
valid_accuracy_list = []

for epoch in range(num_epochs):
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in train_loader:
        images = images.to(device)
        labels = labels.to(device)

        optimizer.zero_grad()

        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    train_loss = running_loss / len(train_loader)
    train_accuracy = correct / total
```

```

# Validation loop
model.eval()
running_loss = 0.0
correct = 0
total = 0

with torch.no_grad():
    for images, labels in valid_loader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

valid_loss = running_loss / len(valid_loader)
valid_accuracy = correct / total

# Store loss and accuracy values
train_loss_list.append(train_loss)
valid_loss_list.append(valid_loss)
train_accuracy_list.append(train_accuracy)
valid_accuracy_list.append(valid_accuracy)

# Print the training/validation statistics
print(f"Epoch: {epoch+1}/{num_epochs} | "
      f"Train Loss: {train_loss:.4f} | Train Acc: {train_accuracy:.4f} | "
      f"Valid Loss: {valid_loss:.4f} | Valid Acc: {valid_accuracy:.4f}")

# Update the learning rate
scheduler.step()

# Save the trained model
torch.save(model.state_dict(), model_name+".pth")

```

```

Epoch: 1/10 | Train Loss: 4.4063 | Train Acc: 0.0498 | Valid Loss: 3.9686 |
Valid Acc: 0.0940
Epoch: 2/10 | Train Loss: 3.3651 | Train Acc: 0.1976 | Valid Loss: 2.7117 |
Valid Acc: 0.3160
Epoch: 3/10 | Train Loss: 2.2224 | Train Acc: 0.4208 | Valid Loss: 1.7309 |
Valid Acc: 0.5420

```


Epoch: 4/10 | Train Loss: 1.5948 | Train Acc: 0.5765 | Valid Loss: 1.4149 | Valid Acc: 0.6200
 Epoch: 5/10 | Train Loss: 1.3335 | Train Acc: 0.6354 | Valid Loss: 1.2661 | Valid Acc: 0.6420
 Epoch: 6/10 | Train Loss: 0.8024 | Train Acc: 0.7861 | Valid Loss: 0.7375 | Valid Acc: 0.8120
 Epoch: 7/10 | Train Loss: 0.6867 | Train Acc: 0.8154 | Valid Loss: 0.8072 | Valid Acc: 0.7540
 Epoch: 8/10 | Train Loss: 0.6515 | Train Acc: 0.8259 | Valid Loss: 0.6773 | Valid Acc: 0.8000
 Epoch: 9/10 | Train Loss: 0.6279 | Train Acc: 0.8351 | Valid Loss: 0.5994 | Valid Acc: 0.8240
 Epoch: 10/10 | Train Loss: 0.5808 | Train Acc: 0.8446 | Valid Loss: 0.6475 | Valid Acc: 0.8180

```
[19]: import matplotlib.pyplot as plt
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
# Plot training and validation loss
plt.plot(range(1, num_epochs+1), train_loss_list, label='Train Loss')
plt.plot(range(1, num_epochs+1), valid_loss_list, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()

plt.subplot(1, 2, 2)
# Plot training and validation accuracy
plt.plot(range(1, num_epochs+1), train_accuracy_list, label='Train Accuracy')
plt.plot(range(1, num_epochs+1), valid_accuracy_list, label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('Training and Validation Accuracy')
plt.legend()

plt.suptitle(model_name)
plt.show()
```

ViT-8-15-nov-23



```
[20]: # Initialize lists to store true labels and predicted labels
true_labels = []
predicted_labels = []

# Validation loop
model.eval()
with torch.no_grad():
    for images, labels in valid_loader:
        images = images.to(device)
        labels = labels.to(device)

        outputs = model(images)
        loss = criterion(outputs, labels)

        running_loss += loss.item()
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

    # Append true labels and predicted labels
    true_labels.extend(labels.cpu().numpy())
    predicted_labels.extend(predicted.cpu().numpy())

valid_loss = running_loss / len(valid_loader)
valid_accuracy = correct / total
```



```
print(f"Epoch: {epoch+1}/{num_epochs} | "
      f"Train Loss: {train_loss:.4f} | Train Acc: {train_accuracy:.4f} | "
      f"Valid Loss: {valid_loss:.4f} | Valid Acc: {valid_accuracy:.4f}")
```

Epoch: 10/10 | Train Loss: 0.5808 | Train Acc: 0.8446 | Valid Loss: 1.3321 |
Valid Acc: 0.8170

```
[21]: from sklearn.metrics import classification_report

# Generate classification report
classification_rep = classification_report(true_labels, predicted_labels)
print("Classification Report:")
print(classification_rep)
```

Classification Report:

	precision	recall	f1-score	support
0	0.80	0.80	0.80	5
1	0.83	1.00	0.91	5
2	0.75	0.60	0.67	5
3	1.00	1.00	1.00	5
4	1.00	0.80	0.89	5
5	1.00	0.80	0.89	5
6	1.00	0.80	0.89	5
7	1.00	0.60	0.75	5
8	0.67	0.80	0.73	5
9	0.71	1.00	0.83	5
10	1.00	0.80	0.89	5
11	1.00	0.80	0.89	5
12	0.83	1.00	0.91	5
13	0.75	0.60	0.67	5
14	0.83	1.00	0.91	5
15	0.83	1.00	0.91	5
16	0.71	1.00	0.83	5
17	0.62	1.00	0.77	5
18	1.00	1.00	1.00	5
19	0.60	0.60	0.60	5
20	0.80	0.80	0.80	5
21	1.00	0.40	0.57	5
22	0.60	0.60	0.60	5
23	1.00	0.80	0.89	5
24	0.75	0.60	0.67	5
25	1.00	0.80	0.89	5
26	1.00	1.00	1.00	5
27	0.57	0.80	0.67	5
28	0.75	0.60	0.67	5
29	0.80	0.80	0.80	5
30	1.00	1.00	1.00	5

31	0.42	1.00	0.59	5
32	1.00	1.00	1.00	5
33	0.75	0.60	0.67	5
34	0.80	0.80	0.80	5
35	1.00	0.80	0.89	5
36	0.57	0.80	0.67	5
37	0.60	0.60	0.60	5
38	0.83	1.00	0.91	5
39	1.00	0.60	0.75	5
40	0.71	1.00	0.83	5
41	1.00	0.80	0.89	5
42	0.83	1.00	0.91	5
43	0.71	1.00	0.83	5
44	0.67	0.80	0.73	5
45	0.80	0.80	0.80	5
46	1.00	1.00	1.00	5
47	1.00	1.00	1.00	5
48	0.83	1.00	0.91	5
49	1.00	1.00	1.00	5
50	0.83	1.00	0.91	5
51	0.50	0.40	0.44	5
52	1.00	0.60	0.75	5
53	0.80	0.80	0.80	5
54	1.00	0.80	0.89	5
55	0.67	0.80	0.73	5
56	1.00	0.80	0.89	5
57	1.00	0.80	0.89	5
58	1.00	1.00	1.00	5
59	1.00	0.60	0.75	5
60	0.75	0.60	0.67	5
61	1.00	0.80	0.89	5
62	0.83	1.00	0.91	5
63	0.60	0.60	0.60	5
64	0.80	0.80	0.80	5
65	0.50	0.80	0.62	5
66	0.80	0.80	0.80	5
67	1.00	1.00	1.00	5
68	0.71	1.00	0.83	5
69	0.80	0.80	0.80	5
70	1.00	0.60	0.75	5
71	0.83	1.00	0.91	5
72	0.83	1.00	0.91	5
73	0.80	0.80	0.80	5
74	1.00	0.80	0.89	5
75	0.80	0.80	0.80	5
76	0.67	0.80	0.73	5
77	0.80	0.80	0.80	5
78	1.00	0.80	0.89	5

79	0.60	0.60	0.60	5	
80	1.00	0.80	0.89	5	
81	0.71	1.00	0.83	5	
82	1.00	1.00	1.00	5	
83	1.00	0.80	0.89	5	
84	0.83	1.00	0.91	5	
85	0.80	0.80	0.80	5	
86	0.62	1.00	0.77	5	
87	1.00	0.80	0.89	5	
88	0.67	0.80	0.73	5	
89	0.75	0.60	0.67	5	
90	1.00	0.80	0.89	5	
91	1.00	0.80	0.89	5	
92	1.00	0.40	0.57	5	
93	0.75	0.60	0.67	5	
94	1.00	0.80	0.89	5	
95	0.80	0.80	0.80	5	
96	0.83	1.00	0.91	5	
97	1.00	0.60	0.75	5	
98	1.00	0.80	0.89	5	
99	1.00	0.80	0.89	5	
accuracy				0.82	500
macro avg		0.84	0.82	0.82	500
weighted avg		0.84	0.82	0.82	500

```
[22]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

plt.figure(figsize=(8, 5))

# Generate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)
cm_display = ConfusionMatrixDisplay(confusion_matrix=cm,
    ↪display_labels=valid_dataset.classes)
cm_display.plot()

plt.title(model_name)

plt.show()
```

<Figure size 800x500 with 0 Axes>

